

# An Industrial Robotic Knowledge Representation for Kit Building Applications

Stephen Balakirsky, Zeid Kootbally, Craig Schlenoff, Thomas Kramer, and Satyandra Gupta

**Abstract**—The IEEE RAS Ontologies for Robotics and Automation Working Group is dedicated to developing a methodology for knowledge representation and reasoning in robotics and automation. As part of this working group, the Industrial Robots sub-group is tasked with studying industrial applications of the ontology. One of the first areas of interest for this subgroup is the area of kit building or kitting which is a process that brings parts that will be used in assembly operations together in a kit and then moves the kit to the assembly area where the parts are used in the final assembly. This paper examines the knowledge representations that have been developed and implemented for the kitting problem.

## I. INTRODUCTION

Kitting is the process in which several different, but related items are placed into a container and supplied together as a single unit. Kitting itself may be viewed as a specialization of the general bin-picking problem. In industrial assembly of manufactured products, kitting is often performed prior to final assembly. Manufacturers utilize kitting due to its ability to provide cost savings [3] including saving manufacturing or assembly space [13], reducing assembly workers walking and searching times [15], and increasing line flexibility [2] and balance [9].

Several different techniques are used to create kits. A kitting operation where a kit box is stationary until filled at a single kitting workstation is referred to as *batch kitting*. In *zone kitting*, the kit moves while being filled and will pass through one or more zones before it is completed. This paper focuses on batch kitting processes.

In batch kitting, the kit's component parts may be staged in containers positioned in the workstation or may arrive on a conveyor. Component parts may be fixtured, for example placed in compartments on trays, or may be in random orientations, for example placed in a large bin. In addition to the kit's component parts, the workstation usually contains a storage area for empty kit boxes as well as completed kits.

Kitting has not yet been automated in many industries where automation may be feasible. Consequently, the cost of building kits is higher than it could be. We are addressing this problem by building models of the knowledge that will

be required to operate an automated kitting workstation in an agile manufacturing environment. This workstation must be able to cope with variations in kit contents, kit layout, and component supply. Our models include representations for non-executable information about the workstation such as information about parts and kit designs, models of executable information such as actions, preconditions, and effects, and models of the process plan necessary for kit construction. A discussion of the functional requirements for the process plan may be found in [1]. For our automated kitting workstation, we assume that a robot performs a series of pick-and-place operations in order to construct the kit. These operations include:

- 1) Pick empty kit and place on work table.
- 2) Pick multiple component parts and place in kit.
- 3) Pick completed kit and place in full kit storage area.

Each of these actions may be a compound action that includes other actions such as end-of-arm tool changes, path planning, and obstacle avoidance.

It should be noted that multiple kits may be built simultaneously. Finished kits are moved to the assembly floor where components are picked from the kit for use in the assembly procedure. The kits are normally designed to facilitate component picking in the correct sequence for assembly. Component orientation may be constrained by the kit design in order to ease the pick-to-assembly process. Empty kits are returned to the kit building area for reuse.

Although the knowledge requirements described in the previous paragraph have been identified for the kitting domain, they are clearly applicable to many types of industrial robot applications. As such, we expect that these knowledge requirements will serve as the basis for the industrial robot ontology being developed in the IEEE RAS Ontologies for Robotics and Automation Working Group [11] (henceforth referred to as the IEEE WG). Throughout the process of developing the kitting ontology, the group will constantly look at the applicability of the requirements outside of kitting and move the pertinent knowledge “up” the ontology (whether in the portion that models the kitting sub-domain, the industrial robot domain, or the upper ontology), as appropriate.

In keeping with the standards philosophy of the IEEE WG, we require the models being developed to be as widely applicable as possible. Therefore, we have created a layered model abstraction where users may adopt as many of the layers of the abstraction as make sense for their specific application. The architecture shown in Figure 1, though developed for the kitting ontology, can be equally applicable to the implementation of any type of formal manufacturing

S. Balakirsky and C. Schlenoff are with the Intelligent Systems Division, National Institute of Standards and Technology, Gaithersburg, MD, USA (e-mail: stephen.balakirsky@nist.gov, craig.schlenoff@nist.gov)

Z. Kootbally is with the Department of Mechanical Engineering, University of Maryland, College Park, MD, USA (email: zeid.kootbally@nist.gov)

T. Kramer is with the Department of Mechanical Engineering, Catholic University of America, Washington, DC, USA (email: thomas.kramer@nist.gov)

S. Gupta is with the Maryland Robotics Center, University of Maryland, College Park, MD, USA (email: skgupta@umd.edu)

knowledge representation. Said in a different way, the implementor can plug in a knowledge representation for a different domain and the architecture would still be valid. In a similar manner, different planning language abstractions could be utilized in the planning language layer of the abstraction and different planning/execution systems could be utilized in the robot language layer of the abstraction.

Specifics on the architecture may be found in Section II. An example of the knowledge representations and the flow from one to the next is presented in Section III. Finally, conclusions and future work may be found in Section IV.

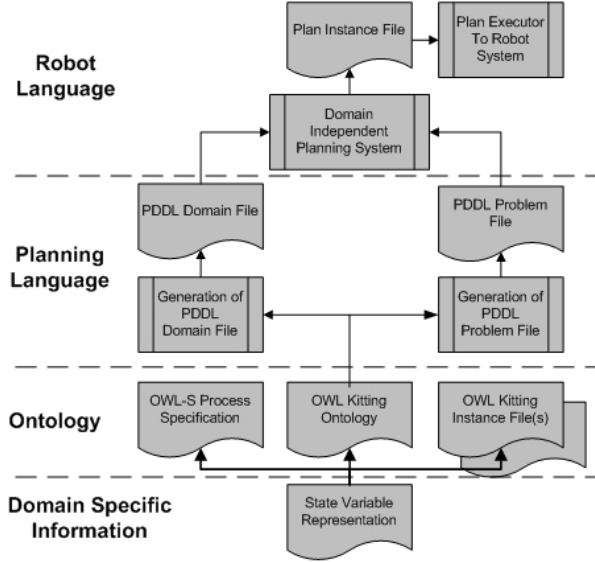


Fig. 1. Kitting data flow abstraction.

## II. ARCHITECTURE DESCRIPTION

The main focus of this work is on the development of knowledge models that allow a kitting workstation to construct kits in an agile manufacturing environment. However, in order to validate these knowledge models, we felt that it was important to be able to utilize the models to construct kit building plans, and then to execute these plans in dynamic virtual and real environments. Due to the advent of open source robotic operating systems such as ROS<sup>1</sup>[6] and simulation packages such as USARSim [4], we do not need to design these systems ourselves. However, our architecture must be designed to represent the required knowledge base in several different abstractions that are likely to be required by these systems as the knowledge flows from domain and process specification, to plan generation, to plan execution. As shown in Figure 1, the abstraction is decomposed into four distinct layers of *Domain Specific Information*, *Ontology*, *Planning Language*, and *Robot Language* that correspond to these knowledge requirements. In this paper,

<sup>1</sup>Certain commercial/open source software and tools are identified in this paper in order to explain our research. Such identification does not imply recommendation or endorsement by the authors, nor does it imply that the software tools identified are necessarily the best available for the purpose.

we will discuss the layers of the Domain Specific Information through Planning Language. Implementors of the abstraction are free to connect to the knowledge interface at the layer that makes sense for their particular application. For our simulated kitting workstation, we intend to fully design the Domain Specific Information and the Ontology and then utilize open source tools that will automatically generate the remaining layers and provide a workstation simulation.

### A. Domain Specific Information

The foundation for the knowledge representation is domain specific information that is produced by an expert in the particular field of study. This includes information on items ranging from what actions and attributes are relevant, to what the necessary conditions are for an action to occur and what the likely results of the action are. We have chosen to encode this basic information in a formalism known as a state variable representation (SVR) [14]. This information will then flow up the abstraction and be transformed into the ontology, planning language, and robot language. In a SVR, each state is represented by a tuple of values of  $n$  state variables  $\{x_1, \dots, x_n\}$ , and each action is represented by a partial function that maps this tuple into some other tuple of values of the  $n$  state variables.

To build the SVR, the group has taken a very systematic approach of identifying and modeling the concepts. Because the industrial robot field is so broad, the group decided to limit its efforts to a single type of operation, namely kitting. A scenario was developed that described, in detail, the types of operations that would be performed in kitting, the sequencing of steps, the parts and machines that were needed, constraints on the process such as pre- and post-conditions, etc. For this scenario, a set of concepts were extracted and defined. These concepts served as the initial requirements for the kitting SVR. The concepts were then modeling in our SVR, building off of the definitions and relationships that were identified in the scenario.

A SVR relies on the elements of constant variable symbols, object variable symbols, state variable symbols, rigid relations, and planning operators. These are defined for the kitting domain in the rest of this section.

1) *Constant Variable Symbols*: For the kitting domain, there is a finite set of constant variable symbols that must be represented. In the SVR, constant variable symbols are partitioned into disjoint classes corresponding to the objects of the domain. The finite set of all constant variable symbols in the kitting domain is partitioned into the following sets:

- A set of *Part*: *Parts* are the basic items that will be used to fill a kit.
- A set of *PartsTray*: *Parts* arrive at the workstation in *PartsTrays*. Each *Part* is at a known position in the *PartsTray*. Each *PartsTray* contains one type of *Part*.
- A set of *KitTray*: A *KitTray* can hold *Parts* in known positions.
- A set of *Kit*: A *Kit* consists of a *KitTray* and, possibly, some *Parts*. A *Kit* is empty when it does not contain

any *Part* and finished when it contains all the *Parts* that constitute a kit.

- A symbol *WorkTable*: A *WorkTable* is an area in the kitting workstation where *KitTrays* are placed to build *Kits*.
- A set of *LargeBoxWithKits*: A *LargeBoxWithKits* contains only finished *Kits*.
- A set of *LargeBoxWithEmptyKitTrays*: A *LargeBoxWithEmptyKitTrays* is a box that contains only empty *KitTrays*.
- A set of *Robot*: A *Robot* in the kitting workstation is a robotic arm that can move objects in order to build *Kits*.
- A set of *VacuumEffectorSingleCup*: *VacuumEffectorSingleCups* are used in a kitting workstation to manipulate *Parts*, *PartsTrays*, *KitTrays*, and *Kits*. A *VacuumEffectorSingleCup* is attached to a *Robot*.
- A set of *EndEffectorHolder*: An *EndEffectorHolder* is a storage unit that holds one type of *EndEffector*.

2) *Object Variable Symbols*: Object variable symbols are typed variables which range over a class or the union of classes of constant variable symbols. Examples of object variable symbols are  $r \in \text{Robot}$ ,  $kt \in \text{KitTray}$ , etc.

3) *State Variable Symbols*: A state variable  $x$  symbol is defined as follows:  $x : A_1 \times \dots \times A_i \times S \rightarrow B_1 \cup \dots \cup B_j$  ( $i, j \geq 1$ ) is a function from the set of states ( $S$ ) and at least one set of constant variable symbols  $A_1 \times \dots \times A_i$  into a set of constant variable symbols  $B_1 \cup \dots \cup B_j$ .

The use of state variable symbols reduces the possibility of inconsistent states and generates a smaller state space. The following state variable symbols are used in the kitting domain:

- *effloc*:  $\text{VacuumEffectorSingleCup} \times S \rightarrow \text{Robot} \cup \text{EndEffectorHolder}$  designates the location of a *VacuumEffectorSingleCup* in the workstation: Attached to a *Robot* or placed in an *EndEffectorHolder*.
- *reff*:  $\text{Robot} \times S \rightarrow \text{VacuumEffectorSingleCup} \cup \{\text{nil}\}$  designates the *VacuumEffectorSingleCup* attached to the *Robot* if there is one attached, otherwise *nil*.
- *onworktable*:  $\text{WorkTable} \times S \rightarrow \text{Kit} \cup \text{KitTray} \cup \{\text{nil}\}$  designates the object placed on the *WorkTable*: A *Kit*, a *KitTray*, or nothing.
- *kitloc*:  $\text{Kit} \times S \rightarrow \text{LargeBoxWithKits} \cup \text{WorkTable} \cup \text{Robot}$  designates the different possible locations of a *Kit* in the workstation: in a *LargeBoxWithKits*, on the *WorkTable*, or being held by a *Robot*.
- *kittrayloc*:  $\text{KitTray} \times S \rightarrow \text{LargeBoxWithEmptyKitTrays} \cup \text{Robot} \cup \text{WorkTable}$  designates the different possible locations of a *KitTray* in the workstation: In a *LargeBoxWithEmptyKitTrays*, being held by on a *Robot*, or on a *WorkTable*.
- *partloc*:  $\text{Part} \times S \rightarrow \text{PartsTray} \cup \text{Kit} \cup \text{Robot}$  designates the different possible locations of a *Part* in the workstation: In a *PartsTray*, in a *Kit*, or being held by a *Robot*.
- *rhold*:  $\text{Robot} \times S \rightarrow \text{KitTray} \cup \text{Kit} \cup \text{Part} \cup \{\text{nil}\}$ : designates the object being held by a *Robot*: A *KitTray*, a

*Kit*, a *Part*, or nothing.

- *islbwfull*:  $\text{LargeBoxWithKits} \times S \rightarrow \{0\} \cup \{1\}$  designates if a *LargeBoxWithKits* is full or not.
- *islbwempty*:  $\text{LargeBoxWithEmptyKitTrays} \times S \rightarrow \{0\} \cup \{1\}$  designates if a *LargeBoxWithEmptyKitTrays* is empty (1) or not (0).
- *ispartstrayempty*:  $\text{PartsTray} \times S \rightarrow \{0\} \cup \{1\}$  designates if a *PartsTray* is empty or not.
- *efftype*:  $\text{VacuumEffectorSingleCup} \times S \rightarrow \text{KitTray} \cup \text{Kit} \cup \text{Part}$  designates the type of object a *VacuumEffectorSingleCup* can hold: A *KitTray*, a *Kit*, or a *Part*.
- *eeffhtype*:  $\text{EndEffectorHolder} \times S \rightarrow \text{VacuumEffectorSingleCup} \cup \{\text{nil}\}$  designates the object an *EndEffectorHolder* is holding: A *VacuumEffectorSingleCup* or nothing.

4) *Sets of Literals*: In order to generate PDDL files, each state variable symbol must be converted into a set of literals (i.e., atoms and negation of atoms). For instance, in order to use the state variable symbol *partloc*:  $\text{Part} \times S \rightarrow \text{PartsTray} \cup \text{Kit} \cup \text{Robot}$  in PDDL, this state variable symbol is converted into the set of literals  $\{\text{partloc}(p,pt), \text{partloc}(p,k), \text{partloc}(p,r)\}$ , where  $p \in \text{Part}$ ,  $k \in \text{Kit}$ ,  $k \in \text{KitTray}$ , and  $r \in \text{Robot}$ . The representation of the location of a *Part* in PDDL is performed by using the correct literal within this set of literals.

5) *Planning Operators*: A planning operator is used to describe one of the procedures that can be used to move from one state to another. A planning operator [14] is a triple  $o=(\text{name}(o), \text{precond}(o), \text{effects}(o))$  whose elements are as follows:

- *name(o)*, the name of the operator, is a syntactic expression of the form  $n(x_1, \dots, x_k)$ , where  $n$  is a symbol called an operator symbol,  $x_1, \dots, x_k$  are all of the object variable symbols that appear anywhere in  $o$ , and  $n$  is unique (i.e., no two operators can have the same operator symbol).
- *precond(o)* and *effects(o)*, the preconditions and effects of  $o$ , respectively, are sets of literals.

The kitting domain presented in this paper is composed of eight operators: *take-kit-tray*, *put-kit-tray*, *take-kit*, *put-kit*, *take-part*, *put-part*, *attach-eff*, and *remove-eff*. Due to space restriction, this section describes only one of these operators as follows:

- *take-kit-tray*( $r, kt, lbwekt, eff, wtable$ ): The *Robot*  $r$  equipped with the *VacuumEffectorSingleCup*  $eff$  picks up the *KitTray*  $kt$  from the *LargeBoxWithEmptyKitTrays*  $lbwekt$ . The *WorkTable*  $wtable$  must be *a priori* empty.

<i>precond</i>	<i>effects</i>
$\text{rhold-empty}(r),$	$\neg \text{rhold-empty}(r),$
$\text{lbwekt-not-empty}(lbwekt),$	$\text{kittrayloc}(kt, r),$
$\text{r-with-eff}(r, eff),$	$\text{rhold}(r, kt),$
$\text{kittrayloc}(kt, lbwekt),$	$\neg \text{kittrayloc}(kt, lbwekt)$
$\text{effloc}(eff, r),$	
$\text{worktable-empty}(wtable),$	
$\text{efftype}(eff, kt)$	

Note that an action  $a$  can be obtained by substituting

the object variable symbols that appear anywhere in the operator with constant variable symbols.

## B. Ontology

Knowledge models may take many forms ranging from informal natural language, to XML schemas, to ontologies. For the development of the knowledge representation, the industrial robots sub-group has decided to use the Web Ontology Language (OWL) [16] as the knowledge representation language. OWL is a family of knowledge representation languages for authoring ontologies and is endorsed by the World Wide Web Consortium (W3C). OWL was chosen by the group because of its popularity among the ontology development community, its endorsement by the W3C, as well as the number of tools and reasoning engines that are available. OWL was also selected as the representation language that will be used in the overall IEEE WG efforts.

In addition to OWL, the industrial robots subgroup will also be using OWL-S [12] to represent the processes and actions that the robot will perform. OWL-S is an ontology built on top of OWL by the DARPA Agent Markup Language (DAML) program [5] for describing Semantic Web Services. Many of the constructs that are used to describe services are equally applicable to encoding our SVR. For example, concepts such as preconditions, results, inputs, outputs, effects, and participants are generic enough to be applied to just about any type of process specification.

Since this work is being directed at the IEEE RAS Ontologies for Robotics and Automation Working Group, it is appropriate that our domain specific knowledge be encoded in an OWL ontology at the first domain independent layer of our abstraction. The knowledge contained in this layer is derived from our SVR and may also contain information that has been collected over other domains.

As more detailed scenarios are determined and a richer set of concepts are uncovered, the ontology will be partitioned based upon the generality of the concept, with the most generally applicable concepts being “higher” in the ontology so they are available to other domains and the more detailed concepts being “lower” because they will likely be very specific to the kitting area. An example of a general concept may be a *Robot* while a specific concept may be a *KitTray*.

As shown in Figure 1, the information in our ontology is divided into three sets of files and consists of a representation of the process specification, the kitting ontology, and the instance file. The process specification file is based off of the planning operators from the SVR and contains descriptions of the individual actions and sequences necessary to construct a kit, e.g. gripping a component from a tray. The kitting ontology file is based off of the constant variable symbols and state variable symbols and contains the concepts related to the specific items that compose the kitting domain, e.g. the hierarchy of what it means to be a *part\_a*. That a *part\_a* is a type of *Part*, and that all *Parts* contain properties such as the part’s weight, dimensions, and grip points. The instance file is based off of the SVR object variable symbols and contains specific information on *this* particular kitting

problem and configuration, e.g. *KitTray kit\_tray\_1* contains four *Parts* of type *part\_a*.

While this file set provides a complete description of the problem domain and environment, most planning systems cannot directly ingest information from an ontology. Therefore, the next layer of the data abstraction known as the Planning Language layer was created.

## C. Planning Language

The Planning Domain Definition Language (PDDL) [7] is an attempt by the domain independent planning community to formulate a standard language for planning. A community of planning researchers has been producing planning systems that comply with this formalism since the first International Planning Competition held in 1998. This competition series continues today, with the seventh competition being held in 2011. PDDL is constantly adding extensions to the base language in order to represent more expressive problem domains. Our work is based on PDDL Version 3.0.

By placing our knowledge in a PDDL representation, we enable the use of an entire family of open source planning systems. As shown in Figure 1, each PDDL file-set consists of two files that specify the domain and the problem. The PDDL domain file is composed of four sections that include requirements, types and constants, predicates, and actions. This file may be automatically generated from a combination of information that is contained in the OWL-S process specification file and the OWL Kitting Ontology file. This auto-generation uses a modified version of the algorithms from Klusch and Gerber [10] that directly generate PDDL syntax.

The requirements section specifies which extensions this problem domain relies on. The planning system can examine this statement to determine if it is capable of solving problems in this domain. In PDDL, all variables that are used in the domain must be typed. Types are defined in the types section. It is also possible to have constants that specify that all problems will share this single value. For example, in the simplest kitting workstation we will have a single *Robot robot\_1*. Predicates specify relationships between instances. For example, an instance of a *KitTray*, *kit\_tray\_1*, can have a physical location and contains instances of *Parts*, *part\_a\_1*, *part\_a\_2*, *part\_b\_1*, and *part\_c\_1*. The final section of the PDDL domain file is concerned with actions. An action statement specifies a way that a planner affects the state of the world. The statement includes parameters, preconditions, and effects. The preconditions dictate items that must be initially true for the action to be legal. The effect equation dictates the changes in the world that will occur due to the execution of the action.

The second file of the PDDL file-set is a problem file. The problem file specifies information about the specific instance of the given problem. This file contains the initial conditions and definition of the world (in the init section) and the goal state that the world must be brought to (in the goal section). A specific example of the ontology to planning language conversion is provided in Section III.

### III. EXAMPLE OF OPERATION

The purpose of this section is to illustrate the various knowledge representations, depicted in Figure 1, and the flow from one knowledge representation to the next.

#### A. State Variable Representation

In the example, the *Robot* has to build a kit that contains two *Parts* of type A, one *Part* of type B and one *Part* of type C. The kitting process is completed once the kit is placed in the *LargeBoxWithKits*. Section III-B details the steps that build the PDDL domain file while Section III-C discusses the process that builds the PDDL problem file.

The state variable representation defines the necessary parameters for the PDDL domain and problem files. The PDDL domain file contains predicates, characterized by state variable symbols, and actions. The PDDL problem file contains the objects present in the problem instance, characterized by constant variable symbols, the initial state  $s_0$  description and the goal state  $s_G$ .

a) *Constant Variable Symbols*: The kitting domain proposed for the example contains a *Robot* *robot\_1*, a *KitTray* *kit\_tray\_1*, a *LargeBoxWithEmptyKitTrays* *empty\_kit\_tray\_supply*, a *LargeBoxWithKits* *finished\_kit\_receiver*, a *WorkTable* *work\_table\_1*, three *PartsTrays* *part\_a\_tray*, *part\_b\_tray*, and *part\_c\_tray*, four *Parts* *part\_a\_1*, *part\_a\_2*, *part\_b\_1*, and *part\_c\_1*, two *VacuumEffectorSingleCup* *part\_gripper* and *tray\_gripper*, and two *EndEffectorHolders* *part\_gripper\_holder* and *tray\_gripper\_holder*. Since a *Kit* is by definition a *KitTray* that contains *Parts*, the kitting domain also contains a constant variable symbol *kit\_1* from *Kit*.

b) *State Variable Symbols*: The state variable symbols for this example are the ones defined in Section II-A.3.

c) *Initial State*: The initial state  $s_0$  can be represented by the set of state variable symbols depicted in Table I.

TABLE I  
INITIAL STATE  $s_0$

```
r-no-eff(robot_1)
kitrayloc(robot_1,empty_kit_tray_supply)
lbwekt-not-empty(empty_kit_tray_supply)
partloc(part_a_1,part_a_tray)
lbwek-not-full(finished_kit_receiver)
partloc(part_a_2,part_a_tray)
part-tray-not-empty(part_a_tray)
partloc(part_b_1,part_b_tray)
part-tray-not-empty(part_b_tray)
partloc(part_c_1,part_c_tray)
part-tray-not-empty(part_c_tray)
efftype(part_gripper,part_a_1)
eff-location(part_gripper,part_gripper_holder)
efftype(part_gripper,part_a_2)
effloc(tray_gripper,tray_gripper_holder)
efftype(part_gripper,part_b_1)
effhold-eff(part_gripper_holder,part_gripper)
efftype(part_gripper,part_c_1)
effhold-eff(tray_gripper_holder,tray_gripper)
efftype(tray_gripper,kit_tray_1)
worktable-empty(work_table_1)
efftype(tray_gripper,kit_1)
```

d) *Goal State*: The goal state  $s_G$  can be represented by the set of state variable symbols depicted in Table II.

TABLE II  
GOAL STATE  $s_G$

```
partloc(part_a_1,kit_1)
partloc(part_a_2,kit_1)
partloc(part_b_1,kit_1)
partloc(part_c_1,kit_1)
kitloc(kit_1,finished_kit_receiver)
```

#### B. Process for the PDDL Domain File

Once the state variable representation is defined, an OWL-S process specification and OWL ontology is built. This must be constructed by hand from the information in the SVR as described in Section II-B. For this example, the specific action *take-kit-tray*'s encoding into OWL-S will be discussed.

1) *OWL-S Representation*: Once the action is specified in the SVR, we can use that information to create an OWL-S process. We are modeling the action as a process and including information about its data inputs and outputs, the preconditions that have to be true for it to be performed, and the result that will be true after the process is executed. In our example, the process is an atomic process because it only involves a single interaction and consists of only one step. If it were a more complex process that involved multiple sub-actions, it would be modeled by an OWL-S composite process.

In the SVR, the preconditions clearly map to the OWL-S preconditions. The rules point to classes and instances in the ontology that model the concepts of kit tray (*kt*), a set of large box with empty kit trays (*lbwekt*), a robot (*r*), and a robot end effector (*eff*). The SVR effects map to the OWL-S results and are also represented in one of the rules languages above. In the case of the *take-kt* action, the result would specify that the location of the kit tray is no longer in a fixed location and is now in the robot end effector.

Though not explicitly represented in the SVR, data inputs and outputs are an important part of the OWL-S representation and can be inferred from the SVR. Specifically, it needs to know which robot is performing the action (*robot\_1*), which kit tray needs to be picked up (*kit\_tray\_1*), which end effector is on the robot (*tray\_gripper*), and from which box the robot needs to pick up the kit tray (*empty\_kit\_tray\_supply*). The output of this action would be a Boolean stating whether the action was completed successfully or not.

2) *PDDL Domain File*: An external tool is used to generate the PDDL Domain File from the OWL/OWL-S representation. Due to space constraints, only an excerpt of the PDDL domain file is displayed in Figure 2.

#### C. Process for the PDDL Problem File

The PDDL problem file is generated by a chain of data conversions that take the SVR through an OWL representation to the PDDL representation.

1) *OWL Representation*: The OWL instance files shown in Figure 1 are derived from the SVR initial and goal states ( $s_0$ ,  $s_G$ ). These files contain specific instances of the classes

```

(define (domain kitting-domain)
  (:requirements :strips :typing)
  (:types
    Robot
    KitTray
    LargeBoxWithEmptyKitTrays
    ...)
  (:predicates
    (rhold-empty ?r - Robot)
    (kittrayloc ?kt - KitTray ?r - Robot)
    ...)
  (:action take-kit-tray
    :parameters(
      ?r - Robot
      ?kt - KitTray
      ?lbwekt - LargeBoxWithEmptyKitTrays
      ...)
    :precondition(and (
      rhold-empty ?r)
      (lbwekt-not-empty ?lbwekt)
      ...)
    :effect(and (
      rhold ?r ?kt)
      (not (kittrayloc ?kt ?lbwekt))
      ...))
  ...))

```

Fig. 2. Excerpt of the PDDL domain file for kitting.

that have been previously defined in the OWL ontology. Two specific instance files are required for the representation with both files containing the same class structures. The difference between the two is that in the *InitFile* the specific properties of the instantiated classes represent the initial state of the domain, and in the *GoalFile* the specific properties of the instantiated classes represent the goal state of the domain.

```

(define (problem kitting-problem)
  (:domain kitting-domain)
  (:objects
    robot_1 - Robot
    work_table_1 - WorkTable
    part_gripper - VacuumEffectorSingleCup
    part_gripper_holder - EndEffectorHolder
    ...)
  (:init
    (r-no-eff robot_1)
    (effloc part_gripper part_gripper_holder)
    (worktable-empty work_table_1)
    ...)
  (:goal
    (and
      (partloc part_a_1 kit_1)
      (partloc part_b_1 kit_1)
      ...)))

```

Fig. 3. Excerpt of the PDDL problem file for kitting.

2) *PDDL Problem File*: The PDDL file is once again automatically generated. An excerpt of the PDDL problem file is shown in Figure 3. This file contains the required sections of objects, initial conditions, and goal conditions.

3) *PDDL File Processing*: Once the PDDL files have been generated, a wide variety of planners exist that could be used to generate a PDDL solution file. In our particular implementation, we are utilizing a variant of the classical FF system by Hoffmann and Nebel [8].

#### IV. CONCLUSIONS AND FUTURE WORK

In this paper we have presented the first knowledge model that is part of the IEEE RAS Ontologies for Robotics and Automation, Industrial Robotics subgroup on kitting. This abstraction begins with data that has been created by an expert user as a state variable representation and encoded

into an OWL ontology. The knowledge is then automatically transitioned into a Planning Domain Definition Language that may be used by a class of planners to find a solution that will provide for actions that will transition the system for its initial state to a prescribed goal state.

This workstation is still very much a work in progress. The IEEE working group has an active mailing list and meets in person at both the IEEE ICRA and IROS conferences. At present, we are working on connecting the planned action sequence through the ROS control stack to our USARSim virtual world. As the knowledge representation is exercised, it is expected that shortcomings and missing data will be detected. Our representation will be updated as necessary to cope with these developments. In addition, the representation will be expanded to include metrics that allow us to represent not only the the information necessary for the construction of a kit, but also for the representation of the quality of the finished kit and the kit building process.

#### REFERENCES

- [1] S. Balakirsky, Z. Kootbally, T. Kramer, R. Madhavan, C. Schlenoff, and M. Shneier. Functional Requirements of a Model for Kitting Plans. In *Proceedings of the 2012 Performance Metrics for Intelligent Systems (PerMIS'12)*, 2012.
- [2] Y. A. Bozer and L. F. McGinnis. Kitting Versus Line Stocking: A Conceptual Framework and Descriptive Model. *International Journal of Production Economics*, 28:1–19, 1992.
- [3] O. Carlsson and B. Hensvold. Kitting in a High Variation Assembly Line. Master's thesis, Luleå University of Technology, 2008.
- [4] S. Carpin, M. Lewis, Jijun Wang, S. Balakirsky, and C. Scraper. USARSim: A Robot Simulator for Research and Education. In *Robotics and Automation, 2007 IEEE International Conference on*, pages 1400–1405, april 2007.
- [5] DARPA. The DARPA Agent Markup Language Homepage. In <http://www.daml.org>, 2012.
- [6] Willow Garage. Robot Operating System (ROS). In <http://www.willowgarage.com/pages/software/ros-platform>, 2012.
- [7] M. Ghallab, A. Howe, C. Knoblock, D. McDermott, A. Ram, M. Veloso, D. Weld, and D. Wilkins. PDDL—The Planning Domain Definition Language. Technical Report CVC TR98-003/DCS TR-1165, Yale, 1998.
- [8] J. Hoffmann and B. Nebel. The FF Planning System: Fast Plan Generation Through Heuristic Search. *Journal of Artificial Intelligence Review*, 14:253–302, 2001.
- [9] J. Jiao, M. M. Tseng, Q. Ma, and Y. Zou. Generic Bill-of-Materials-and-Operations for High-Variety Production Management. *Concurrent Engineering: Research and Applications*, 8(4):297–321, December 2000.
- [10] M. Klusch and A. Gerber. Semantic web service composition planning with owls-xplan. In *In Proceedings of the 1st Int. AAAI Fall Symposium on Agents and the Semantic Web*, pages 55–62, 2005.
- [11] R. Madhavan, W. Yu, G. Biggs, and C. Schlenoff. IEEE RAS STanding Committee for Standards Activities: History and Status Update. *Journal of Advanced Robotics Special Issue on Internationalization*, 2011.
- [12] D. Martin, M. Burstein, J. Hobbs, O. Lassila, D. McDermott, S. McIlraith, S. Narayanan, M. Paolucci, B. Parsia, T. Payne, E. Sirin, N. Srinivasan, and K. Sycara. OWL-S: Semantic Markup for Web Services. In <http://www.w3.org/Submission/OWL-S/>, 2012.
- [13] L. Medbo. Assembly Work Execution and Materials Kit Functionality in Parallel Flow Assembly Systems. *International Journal of Production Economics Journal of Industrial Ergonomics*, 31:263–281, 2003.
- [14] D. Nau, M. Ghallab, and P. Traverso. *Automated Planning: Theory & Practice*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2004.
- [15] G.F. Schwind. How Storage Systems Keep Kits Moving. *Material Handling Engineering*, 47(12):43–45, 1992.
- [16] W3C. OWL 2 Web Ontology Language Document Overview. In <http://www.w3.org/TR/owl-overview/>, 2012.